

MindfulAI

Group #6

Sumedh Tuttagunta, Andrew Tran, Tuan Do

1: Ethical Business Plan

1.A. Company Name of Fictitious Company

The name of our fictitious technology company is MindfulAI. MindfulAI is an AI-powered mental health support company that provides accessible, private, and evidence-based mental wellness tools for people who may not have immediate access to traditional therapy. The company is designed as a support companion, not as a replacement for licensed therapists, doctors, or crisis services. MindfulAI focuses on helping users reflect on their emotions, practice coping strategies, track mental wellness patterns, and connect with human support when needed.

MindfulAI is built around the belief that technology can help reduce barriers to mental health care, but only when it is designed with strong ethical safeguards. Mental health data is deeply personal, and users seeking help may be vulnerable, stressed, isolated, or in crisis. Because of this, MindfulAI treats safety, privacy, transparency, and human oversight as core business responsibilities rather than optional features. The company's goal is not simply to grow quickly, but to grow responsibly while protecting the people who trust the platform.

1.B. Long-Term Vision Statement

1.B.1 Goals

At MindfulAI, our goal is to provide professional-grade mental wellness support at affordability and scale while still respecting the limits of artificial intelligence. We are developing a global mental health companion powered by AI that can help users manage stress, anxiety, low mood, and daily emotional challenges. The platform will use evidence-based techniques such as cognitive behavioral therapy exercises, guided journaling, emotional check-ins, grounding exercises, breathing exercises, and coping plans. It will also include crisis detection systems that identify when a user may need immediate human support.

Our main goal is to help close the global treatment gap for mental health by giving people a safe first layer of support when traditional therapy is unavailable, unaffordable, delayed, or difficult to access. MindfulAI is also designed to support traditional therapists by helping users practice skills between sessions and monitor their mental wellness over time. The platform can also act as a gateway for people who are afraid of traditional therapy or who are unsure how to begin asking

for help. This goal connects to the global need for expanded mental health support described by the World Health Organization [1].

Within five years, MindfulAI hopes to reach millions of users while maintaining a clear ethical standard. The product must help users build healthier lives, not become dependent on the app. We want to support early intervention for common mental health challenges, reduce stigma around seeking help, and create a safe path toward human support when needed. However, our success will not be measured only by user growth. It will also be measured by clinical safety, privacy protection, user trust, fairness across communities, and responsible escalation to human care.

MindfulAI's long-term goal is to become a trusted mental health technology company that proves AI can be used responsibly in a sensitive area. The company will not claim that AI can replace therapists or medical professionals. Instead, MindfulAI will focus on being a careful, transparent, and ethical support companion that helps users understand their emotions, practice healthy coping strategies, and connect with real human help when necessary.

1.B.2 Idea Origination

The idea for MindfulAI came from the founding team's personal experiences with friends and family members who struggled to find affordable, stigma-free mental health support. Some people were willing to get help but could not afford therapy. Others lived in areas where appointments were unavailable for months. Some were afraid of being judged, misunderstood, or labeled because of their mental health challenges. These experiences made our team realize that access to mental health support is not only a medical issue, but also a social and ethical issue.

As we studied AI ethics and healthcare applications, we noticed that AI-based mental health tools were becoming more common, but many of them raised serious ethical questions. Some products seemed to overstate what AI could safely do. Others did not clearly explain how user data was stored, shared, or used. Some apps focused more on engagement than actual user well-being, which can be especially dangerous in the mental health space. The National Institute of Mental Health also explains that mental health technology has promise, but users need to understand its limits and risks [2].

MindfulAI was created as a response to this problem. We wanted to design a company that does not treat ethics as a separate section added after the product is built. Instead, ethics is part of the business model from the beginning. Our startup idea is based on the belief that AI can help expand access to mental wellness support, but only if the product is transparent about its limits, clinically reviewed, privacy-focused, culturally inclusive, and connected to human care when needed.

MindfulAI is also based on the idea that mental health technology should not simply copy traditional therapy into a chatbot. The product must be designed around safety, privacy, culture, accessibility, and user control. The goal is not to replace therapists, but to create a responsible support tool that can help people take the first step toward better mental wellness.

1.B.3 Purpose/Values/Mission

Our mission is to provide compassionate, accessible, and evidence-based AI-powered mental health support while emphasizing the continued importance of human connection. MindfulAI exists to reduce barriers to mental wellness support while helping users understand when they need a licensed therapist, doctor, crisis counselor, or trusted person in their life.

Our first value is human-centered design. MindfulAI begins with real human needs, not with technology for its own sake. Every major feature must be reviewed from the user's point of view, especially the point of view of vulnerable users who may be anxious, depressed, isolated, or in crisis.

Our second value is radical transparency. Users should know that they are interacting with AI, what the AI can do, what it cannot do, and when a human professional should be involved. MindfulAI will not market itself as a therapist or claim that it can replace professional diagnosis or treatment.

Our third value is privacy as a basic right. Mental health conversations may include trauma, family conflict, addiction, suicidal thoughts, financial struggles, or other deeply personal information. This information must never be treated as ordinary business data. MindfulAI will collect only what is necessary, protect sensitive information, and give users control over their data.

Our fourth value is clinical responsibility. MindfulAI will use evidence-based wellness techniques and involve licensed mental health professionals in reviewing safety, quality, and appropriateness. The company will not make clinical claims without evidence.

Our fifth value is inclusive accessibility. Mental health is experienced differently across cultures, languages, disabilities, and socioeconomic conditions. Because of that, MindfulAI must be tested with diverse communities and adapted carefully instead of assuming one model works for everyone.

1.B.4 Key Questions

The first key question guiding MindfulAI is: How can we use AI to improve mental wellness without creating harm, dependency, or false trust? This question reminds us that mental health technology is not neutral. A chatbot that gives poor advice, misses a crisis, or makes users overly

dependent can cause real harm. MindfulAI must be designed to encourage healthy coping, real-world support, and proper human care instead of endless app use.

The second key question is: How can MindfulAI scale access while still preserving human connection and professional oversight? Scaling a mental health product is not the same as scaling entertainment software. Users may need empathy, privacy, safety, and sometimes urgent help. MindfulAI must therefore operate as a support tool that works alongside therapists, crisis lines, and healthcare providers, rather than pretending that AI alone is enough.

The third key question is: What governance structures will keep MindfulAI ethical as the company grows and faces commercial pressure? Startups often face pressure to grow faster, collect more data, increase engagement, or reduce human review costs. MindfulAI's ethical challenge is to prevent those business pressures from weakening safety and privacy. This is why the company will use an ethics board, independent audits, clear privacy limits, safety reporting, and user-centered metrics.

1.C. Strategy with Ethical Impacts AND Ethical Safeguards

1.C.1.1 OKR 1 Objective and Key Result

Objective: Develop and evaluate a clinically responsible AI-powered mental wellness companion for users experiencing mild to moderate stress, anxiety, or low mood.

Key Results:

1. By mid-2027, launch a fully functioning MindfulAI platform with at least 20 evidence-informed support techniques, including guided journaling, grounding exercises, thought reframing, emotional check-ins, breathing exercises, and coping plans.
2. By the end of 2027, complete an initial pilot study with 500 users and aim for at least 70 percent of users to report improvement in self-reflection, emotional awareness, or coping confidence after eight weeks of use.
3. Require licensed mental health professionals to review sample conversations and rate the AI's responses for empathy, helpfulness, appropriateness, and safety.
4. Achieve an average professional quality rating of at least 4.0 out of 5.0 before a major version is released to the public.

This OKR matters because mental health support must be judged by more than user growth or app engagement. A regular technology company might measure success by how much time users spend in the product, but MindfulAI must measure whether the app actually helps users build healthy coping skills. The company must avoid designing features that keep users attached to the product without improving their well-being. For this reason, the OKR focuses on safety, quality, user outcomes, and professional review.

The objective is not to claim that MindfulAI cures mental illness or replaces therapy. Instead, the goal is to create a safe support layer for users who need help managing everyday emotional challenges. The product will be clearly positioned as a mental wellness companion and not as a licensed therapist, doctor, or emergency service.

1.C.1.2 OKR 1 Metric(s) with Experiment(s)

Metrics:

1. Clinical Effectiveness: MindfulAI will use standard questionnaires at standard time points to measure anxiety, depression, stress, and coping confidence. Outcomes will be recorded at baseline, 4 weeks, and 8 weeks. The goal is to see whether users report meaningful improvement after using the platform.
2. User Engagement: MindfulAI will measure average sessions per week, completion of exercises, and return use. However, engagement will not be treated as success by itself. The goal is healthy engagement, not addictive use.
3. Safety: MindfulAI will monitor the frequency of escalations to human counselors, crisis resources, or emergency support. The platform will aim to identify crisis situations with very high accuracy and avoid giving harmful advice.
4. Quality: Conversations will be rated by mental health professionals based on empathy, helpfulness, appropriateness, and safety. The average rating must be above 4.0 out of 5.0 before major release.

Experiments:

1. Comparison Study with 500 Participants: In this study, one group receives direct access to MindfulAI while another group remains on an 8-week waiting list and receives general mental health resources. Participants' symptoms and coping confidence are measured at baseline, after 8 weeks, and after 4 months. The goal is to determine whether users with access to MindfulAI report greater improvement than those on the waiting list.
2. Different Conversation Style Testing: MindfulAI will test two different conversation styles, one more structured and directive and one more open and exploratory. The company will evaluate which style leads to better user engagement, stronger coping outcomes, and lower dependency risk.
3. Expert Review: Three licensed therapists will review 100 conversations and provide a quality and safety rating for each conversation. Their ratings will be compared with the AI's confidence scores.
4. Long-Term Tracking Study: MindfulAI will conduct a long-term tracking study with 1,000 users for 6 months under real-life use. This will help the company understand whether users continue to benefit over time, whether dependency risks appear, and whether the platform remains safe outside a controlled testing environment.

1.C.1.3 OKR 1 Ethical Impact(s)/Issues(s)

The first ethical issue is the risk of harm from inappropriate advice. A user may share suicidal thoughts, abuse, trauma, or severe distress. If MindfulAI fails to recognize the seriousness of the situation, the user could be harmed. Even if the AI gives advice that sounds supportive, it may still be inappropriate if the user needs immediate professional help. This is one of the most serious risks for any AI mental health product.

The second ethical issue is misleading users about what AI can do. If MindfulAI markets itself as a therapist or presents itself as emotionally equivalent to a human professional, users may place too much trust in it. This could cause users to avoid therapy, delay treatment, or rely on AI during moments when they need real human support.

The third ethical issue is unhealthy dependency. A user may begin to rely on the AI as their main source of emotional connection. This could reduce real-world relationships and support systems rather than strengthening them. MindfulAI must avoid design choices that encourage endless use instead of healthier independence.

The fourth ethical issue is clinical quality. Mental health techniques must be used carefully because advice that works for one person may not work for another. If the system gives generic advice without understanding context, culture, trauma history, or severity, it could create harm.

The fifth ethical issue is unequal performance across groups. Mental health language and emotional expression vary across cultures, languages, ages, and communities. If the system is trained mostly on narrow data, it may misunderstand or underserve some users.

1.C.1.4 OKR 1 Ethical Safeguards

MindfulAI will use multiple crisis detection systems as a core safeguard. The system will combine keyword detection, conversation context, user behavior patterns, and risk scoring to identify potential crisis situations. When risk is detected, MindfulAI will immediately state that it is an AI system and encourage the user to contact emergency services, a crisis hotline, or a trusted person. When available, the platform will offer a direct connection to crisis support. High-risk conversations will be flagged for human review according to the company's safety protocols.

MindfulAI will also clearly explain its limits during onboarding and throughout the product. Users will be told that MindfulAI is not a doctor, therapist, or emergency service. Marketing claims will be reviewed by the ethics board to prevent exaggeration. The company will not describe the product as a replacement for therapy. It will describe the product as a support companion that can help users practice skills and reflect between human interactions.

To prevent unhealthy dependency, MindfulAI will avoid addictive design patterns. The app will not use streaks, manipulative notifications, emotional pressure, or rewards that encourage constant use. Instead, the product will encourage users to practice coping skills outside the app, connect with trusted people, and gradually build independence. If usage patterns suggest possible dependency, the app will gently remind users to seek human connection and professional support.

MindfulAI will also use professional review before major releases. Licensed mental health professionals will review sample conversations and evaluate whether the AI's responses are safe, appropriate, respectful, and useful. Any feature that creates repeated unsafe responses will be delayed until it is corrected.

Finally, MindfulAI will publish a safety report every three months. This report will summarize crisis escalations, safety improvements, fairness testing, and major product changes. The report will protect user privacy while still giving the public meaningful information about how the company manages risk.

1.C.2.1 OKR 2 Objective and Key Result

Objective: Build a secure and privacy-focused platform that protects user data and gives users meaningful control over their personal information.

Key Results:

1. By 2027, ensure that all user conversations and sensitive data are encrypted and protected.
2. Maintain zero major data breaches or unauthorized access incidents.
3. Achieve at least 95 percent positive feedback from users regarding trust and privacy.
4. Allow all users to easily view, download, correct, and delete their personal data.
5. Successfully pass regular third-party security audits with no major unresolved issues.

This objective is central to MindfulAI because mental health data is among the most sensitive information a person can share. Users may discuss trauma, suicidal thoughts, family conflict, workplace stress, substance use, abuse, medical concerns, or personal identity. If this information is exposed, sold, or misused, the consequences could include embarrassment, discrimination, blackmail, job-related harm, insurance-related harm, or loss of trust in mental health services.

MindfulAI's privacy goal is based on the idea that users should not need to sacrifice their dignity in order to receive mental health support. The company will not sell user conversation data to advertisers, employers, insurers, or data brokers. It will not use sensitive mental health conversations for unrelated commercial targeting. It will not collect unnecessary location data, contact lists, or social media information.

The platform will be designed around data minimization, user consent, encryption, and clear explanations. Users must understand what data is collected, why it is collected, how long it is stored, and how they can remove it. Privacy is not only a technical requirement for MindfulAI. It is an ethical promise to users who may be sharing some of the most personal details of their lives.

1.C.2.2 OKR 2 Metric(s) with Experiment(s)

Metrics:

1. Data Security: MindfulAI will track the number of security incidents, unauthorized access events, and data breaches. The goal is zero major incidents.
2. User Trust: MindfulAI will run user surveys to measure whether users feel safe, informed, and in control of their data. The target is at least 95 percent positive feedback.
3. Privacy Control Usage: MindfulAI will measure whether users can successfully use privacy settings, including viewing, downloading, correcting, and deleting their data. The target is at least 70 percent successful use without needing customer support.
4. Audit Compliance: MindfulAI will complete regular external audits and correct all serious findings before expanding the product.
5. System Reliability: MindfulAI will maintain platform uptime of at least 99.9 percent so users can access support when needed.

Experiments:

1. Privacy Dashboard Testing: MindfulAI will test different versions of the privacy dashboard to see which design helps users better understand and use privacy controls. The company will measure whether users can find the delete, download, and consent settings quickly.
2. User Trust Surveys: MindfulAI will conduct regular surveys to measure whether users understand how their data is used. These surveys will help identify confusing language in privacy explanations.
3. Security Testing: MindfulAI will hire independent security professionals to simulate cyberattacks and identify weaknesses in authentication, encryption, data storage, employee permissions, and API access.
4. Privacy Explanation Testing: MindfulAI will test different ways of explaining privacy policies. The goal is to write privacy information in simple language rather than legal language that most users do not understand.
5. Breach Response Drills: MindfulAI will run internal drills to test how quickly the company can detect a possible breach, protect affected systems, notify users, and correct the problem.

These metrics and experiments are important because privacy promises are not enough. The company must prove through testing that its systems actually protect users.

1.C.2.3 OKR 2 Ethical Impact(s)/Issues(s)

The first ethical issue is data misuse. Mental health data could be extremely valuable to advertisers, insurers, employers, or other third parties. A company under financial pressure may be tempted to monetize user data, even if doing so damages user trust. MindfulAI must reject this business model because users seeking mental health support are often in a vulnerable position.

The second ethical issue is lack of informed consent. Many users click “agree” without understanding what data is being collected or how it may be used. If privacy policies are written in confusing language, then consent is not truly meaningful. A mental health platform has a special duty to explain privacy clearly because users may be emotionally distressed when they sign up.

The third ethical issue is loss of user control. If users cannot delete their data, correct their information, or understand what is stored about them, then the platform takes power away from the people it claims to help. MindfulAI must give users meaningful control over their information instead of hiding important settings.

The fourth ethical issue is employee misuse. Even if the company does not sell data, internal employees could access sensitive conversations without a legitimate reason. This would violate user trust and could create serious harm.

The fifth ethical issue is breach harm. If hackers expose user conversations, the damage could be personal, professional, and long-lasting. Mental health data is not like ordinary browsing data. It can reveal intimate details about a person’s private life.

The sixth ethical issue is pressure to scale the business. As MindfulAI grows, there may be pressure to collect more data, use data for advertising, or reduce privacy protections in order to increase revenue. This is why privacy safeguards must be built into the company’s structure from the beginning.

1.C.2.4 OKR 2 Ethical Safeguards

MindfulAI will use encryption for data in transit and at rest. Access to sensitive data will be limited by role, and employees will only be able to access information needed for their work. All internal access will be logged and reviewed. Employees who access user data without authorization will face disciplinary action, including termination.

Users will have a privacy dashboard where they can view, download, correct, and delete their data. The dashboard will use simple language and will not hide important controls behind

confusing menus. MindfulAI will also provide short explanations at the moment data is collected, instead of relying only on a long privacy policy.

MindfulAI will not sell user data. This rule will be written into the company's ethics policy and reviewed by the ethics board. The company will not sell or share mental health conversations with advertisers, employers, insurers, or data brokers.

The company will also conduct annual third-party security audits and publish a summary of privacy protections in its transparency report. If a breach occurs, MindfulAI will notify affected users quickly, explain what happened, explain what information was affected, and describe what corrective action was taken.

MindfulAI will use privacy-by-design principles, meaning that privacy protections will be built into the product from the beginning rather than added later. The company will also avoid collecting unnecessary information such as contact lists, social media accounts, or precise location data unless there is a clear and user-approved safety reason.

1.C.3.1 OKR 3 Objective and Key Result

Objective: Reduce bias and improve inclusive access across cultures, languages, income levels, disabilities, and demographic groups.

Key Results:

1. By 2028, support at least five major languages with culturally reviewed mental health content.
2. Complete fairness testing across major demographic groups before every major release.
3. Create community advisory panels in at least three regions.
4. Reduce measurable performance gaps between demographic groups to less than 5 percent in safety detection, user satisfaction, and professional review quality scores.

This objective is necessary because mental health is shaped by culture, language, family, religion, race, gender, disability, income, and local social conditions. A phrase that signals distress in one culture may not be expressed the same way in another. Some communities may describe depression through physical symptoms, while others may use emotional, family-based, or spiritual language. If MindfulAI is trained mainly on Western, English-speaking, middle-class data, it may misunderstand users from other backgrounds.

MindfulAI's goal is not just to translate the app into different languages. Translation alone is not enough. The company must also adapt examples, coping strategies, explanations, and escalation resources to the communities it serves. For example, users in rural areas may have different access to therapy than users in urban areas. Users with disabilities may need different interface

options. Users from cultures where mental health stigma is high may need more careful language around support and privacy.

This OKR makes inclusion a measurable part of the business strategy rather than a vague promise.

1.C.3.2 OKR 3 Metric(s) with Experiment(s)

Metrics:

1. Performance Consistency: MindfulAI will test whether crisis detection, response quality, and user satisfaction differ by language, region, age group, disability status, or cultural background when users choose to share that information.
2. Language Quality: Native speakers and mental health professionals will review translated content to ensure that it is accurate, respectful, and culturally appropriate.
3. Accessibility: MindfulAI will measure whether users with visual, motor, cognitive, or language-related disabilities can use the app successfully.
4. Community Review: MindfulAI will track how many major features are reviewed by community advisors before launch.
5. Fairness Gaps: MindfulAI will measure whether certain groups receive lower safety ratings or lower quality responses and will treat those gaps as product safety issues.

Experiments:

1. Fairness Evaluation Study: MindfulAI will test the same mental health scenarios across different demographic and cultural contexts to see whether the AI responds with equal safety and quality.
2. Language Adaptation Review: Instead of directly translating English content, MindfulAI will ask local clinicians and community reviewers to evaluate whether examples and coping strategies make sense in their communities.
3. Accessibility Testing: MindfulAI will test screen reader compatibility, larger text, simple-language mode, voice input, and low-bandwidth access.
4. Community Advisory Panels: MindfulAI will invite users, clinicians, advocates, and community leaders from different backgrounds to review new features before launch.
5. Human Review Comparison: MindfulAI will compare AI-only feedback with human-reviewed feedback in diverse user groups to identify where the AI performs poorly and needs improvement.

1.C.3.3 OKR 3 Ethical Impact(s)/Issues(s)

The first ethical issue is biased performance. If MindfulAI understands one group better than another, then the product may help some users while harming or ignoring others. This is

especially dangerous in mental health because misunderstanding a user's distress can lead to unsafe responses.

The second ethical issue is cultural insensitivity. Some coping suggestions may be inappropriate or unrealistic depending on a user's family structure, religion, income, local environment, or cultural background. A suggestion that sounds simple in one context may be impossible or disrespectful in another.

The third ethical issue is digital inequality. People with older phones, limited internet access, disabilities, or low literacy may be excluded if MindfulAI is designed only for highly connected users. A mental health support platform should not only serve users who already have the most resources.

The fourth ethical issue is language harm. Poor translation can make mental health advice confusing or even dangerous. A phrase used to discuss self-harm, trauma, or panic must be translated with care.

The fifth ethical issue is unequal crisis support. If crisis resources are only accurate in some countries or languages, users outside those areas may receive lower-quality help during moments of serious need.

The sixth ethical issue is power imbalance. A technology company may assume it knows what vulnerable communities need without actually listening to them. This can repeat existing inequalities. MindfulAI must avoid treating communities as test subjects and instead include them as partners in product development.

1.C.3.4 OKR 3 Ethical Safeguard

MindfulAI will create a fairness review process before every major release. No major feature will launch until the company tests whether it performs safely across different user groups. If a feature shows a serious performance gap, the release will be delayed until the issue is fixed.

The company will work with local clinicians, translators, disability advocates, and community advisory panels to adapt the product responsibly. This will include reviewing language, examples, crisis resources, and cultural assumptions. MindfulAI will not assume that a mental health approach developed for one population automatically works for all users.

MindfulAI will also support accessibility features such as screen readers, simple-language mode, adjustable text size, voice options, and low-bandwidth use. These features will help make the product usable for people with different abilities, devices, and internet access levels.

MindfulAI will avoid launching in regions where it cannot provide reliable crisis information or culturally appropriate support. Growth will be limited by safety readiness, not only by business

demand. The ethics board will review expansion plans to make sure the company does not enter new markets before it can responsibly serve users there.

The company will publish fairness summaries in its transparency reports while protecting user privacy. These summaries will explain what was tested, what gaps were found, and what actions were taken to improve the system.

2: Cultural Policy

2.A Core Values

MindfulAI's company culture is based on responsibility, humility, transparency, privacy, inclusion, and care. We want to be thought of as a technology company that understands the seriousness of the problem it is trying to solve. We do not want to be known as a company that uses mental health struggles as a way to collect data, increase screen time, or make exaggerated claims about AI. We want users, clinicians, and communities to see MindfulAI as careful, honest, and human-centered.

Responsibility means that every employee understands the possible consequences of building mental health technology. A small product decision can affect a vulnerable person's safety, privacy, or trust. Because of that, employees are expected to question unsafe decisions and raise concerns early.

Humility means that we admit the limits of our technology. MindfulAI employees should never assume that AI knows more than users, therapists, families, or communities. AI can be useful, but it cannot replace human judgment, professional care, or lived experience.

Transparency means that we explain our product clearly and do not hide behind technical language. Privacy means that we treat mental health data with the highest respect. Inclusion means that we design with different communities instead of designing only for one type of user. Care means that business success must never be separated from user well-being.

2.B Motivation

What we love is the possibility of using technology to reduce loneliness, stigma, and barriers to care. We love the idea that someone who is afraid, embarrassed, or unable to access therapy could still receive a safe first layer of support. We love building tools that help people understand their emotions, practice coping strategies, and take the next step toward human support when needed.

We also love working across disciplines. MindfulAI is not only a software company. It combines computer science, psychology, ethics, design, privacy, accessibility, and public health. Our

culture should encourage employees to think beyond code and consider the real people affected by the product.

What we fear is building a product that looks helpful but quietly causes harm. We fear a future where vulnerable people are pushed toward cheap automated systems instead of receiving real human care. We fear a company culture where growth, investor pressure, or engagement metrics become more important than safety. We also fear privacy violations, biased performance, and misleading marketing.

These fears are important because they keep the company honest. MindfulAI's culture must encourage employees to raise concerns early, question unsafe decisions, and protect users even when doing so slows business growth.

2.C Summary

1. Safety
2. Privacy
3. Transparency
4. Inclusion
5. Care

3: Ethics Policy

3.A Core Items

MindfulAI's first core ethics policy item is user safety. The company will not release any feature that creates unreasonable risk for users experiencing mental health distress. All high-risk features must go through clinical review, safety testing, and ethics board review. Crisis detection, escalation language, and emergency resources must be tested regularly. MindfulAI will not claim to diagnose, cure, or replace therapy.

The second core item is truthful representation of AI. MindfulAI will always tell users that they are interacting with AI. The company will not design the chatbot to trick users into believing it is human. The product will clearly explain that the AI can provide wellness support, reflection prompts, and coping exercises, but it cannot replace licensed care. All marketing claims must be based on evidence and reviewed for accuracy.

The third core item is privacy and data protection. MindfulAI will collect only the data needed to provide the service. The company will not sell mental health conversation data to advertisers, employers, insurers, or data brokers. Users will have clear controls to view, download, and delete

their data. Sensitive data will be encrypted, access will be limited, and employee access will be logged.

The fourth core item is fairness and inclusion. MindfulAI will test its product across languages, cultures, and demographic groups. If the system performs worse for certain groups, the company will treat that as a safety problem, not just a technical issue. The company will work with community advisors, clinicians, and accessibility experts before expanding into new populations.

The fifth core item is human oversight. MindfulAI will maintain connections to licensed professionals, crisis resources, and human review systems. The company will not use AI to make high-risk mental health decisions without human involvement. Human review is especially important for crisis situations, safety incidents, and product changes that affect vulnerable users.

The sixth core item is accountability. MindfulAI will publish regular safety, privacy, and fairness reports. The company will maintain an independent ethics board and give that board authority to review major product decisions, marketing claims, data practices, and expansion plans. Employees will be encouraged to report ethical concerns without fear of punishment.

3.B Board

The first person we would invite to MindfulAI's ethics board is Dr. Timnit Gebru. Dr. Gebru is widely known for her work in AI ethics, algorithmic bias, and accountability. She is the founder and executive director of the Distributed Artificial Intelligence Research Institute. She would be valuable to MindfulAI because our company depends on AI systems that must be tested for bias, fairness, and social harm. Her experience would help the company avoid treating fairness as only a technical metric and instead consider broader questions about power, inequality, and who is harmed when AI systems fail. Her AI ethics work makes her important for reviewing bias, fairness, and accountability risks in MindfulAI [5].

The second person we would invite is Dr. John Torous. Dr. Torous is a psychiatrist and clinical informaticist who works in digital mental health and psychiatry. He is connected to the Division of Digital Psychiatry at Beth Israel Deaconess Medical Center, which focuses on digital mental health, AI safety evaluation, and clinical innovation. He would be valuable to MindfulAI because our company operates in the mental health technology space. His experience with mental health apps, clinical evaluation, and digital psychiatry would help the company understand what responsible evidence, safety, and user support should look like. His work in digital psychiatry makes him directly relevant to MindfulAI's mental health technology sector [3].

The third person we would invite is Dr. Anita L. Allen. Dr. Allen is a scholar of privacy law, ethics, bioethics, and philosophy. She is connected to the University of Pennsylvania and is known for her work on privacy and data protection. She would be valuable to MindfulAI because our company handles extremely sensitive mental health information. Her guidance would help

the company build privacy policies that respect user dignity, not just legal minimums. She would also help the company think about consent, user control, and the ethical risks of collecting sensitive personal data. Her privacy and bioethics background makes her useful for reviewing MindfulAI's sensitive mental health data practices [4].

4: YouTube Presentation

YouTube Presentation URL: [Add YouTube link here]

Our YouTube presentation explains the MindfulAI GitHub wiki documentation and walks through the major parts of the project. The presentation begins with the company name, group number, and team member introductions. Then it explains MindfulAI's long-term vision, including the company's goal of providing accessible AI-powered mental wellness support while preserving privacy, safety, and human connection.

The presentation then discusses the three OKRs. The first OKR focuses on clinical safety and effectiveness. The second OKR focuses on privacy, security, and user control. The third OKR focuses on bias reduction and inclusive access. For each OKR, the presentation explains the metrics, experiments, ethical risks, and safeguards.

After that, the presentation summarizes the Cultural Policy by explaining the company's values and the six-word culture summary. The presentation also summarizes one major point from the Ethics Policy, such as truthful representation of AI or privacy protection. Finally, the presentation introduces one ethics board member and explains why that person was selected.

5: References

[1] World Health Organization. 2021. Mental Health Atlas 2020. World Health Organization. Retrieved May 1, 2026 from <https://www.who.int/publications/i/item/9789240036703>

[2] National Institute of Mental Health. n.d. Technology and the Future of Mental Health Treatment. National Institute of Mental Health. Retrieved May 1, 2026 from <https://www.nimh.nih.gov/health/topics/technology-and-the-future-of-mental-health-treatment>

[3] Harvard Catalyst Profiles. n.d. John Torous. Harvard Catalyst Profiles. Retrieved May 1, 2026 from <https://connects.catalyst.harvard.edu/Profiles/profile/1243692>

[4] Electronic Privacy Information Center. n.d. Anita L. Allen. Electronic Privacy Information Center. Retrieved May 1, 2026 from <https://epic.org/people/anita-l-allen/>

[5] Stanford Institute for Human-Centered Artificial Intelligence. 2022. Timnit Gebru: Ethical AI Requires Institutional and Structural Change. Stanford HAI. Retrieved May 1, 2026 from <https://hai.stanford.edu/news/timnit-gebru-ethical-ai-requires-institutional-and-structural-change>